

Practical 1

```
In [ ]: NAME: Gayatri Ghorpade  
ROLLNO:14144
```

```
In [30]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt  
import warnings  
warnings.filterwarnings('ignore')
```

```
In [31]: df=pd.read_csv("uber.csv")  
df
```

Out[31]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	p
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	
...
199995	42598914	2012-10-28 10:49:00.00000053	3.0	2012-10-28 10:49:00 UTC	-73.987042	
199996	16382965	2014-03-14 01:09:00.0000008	7.5	2014-03-14 01:09:00 UTC	-73.984722	
199997	27804658	2009-06-29 00:42:00.00000078	30.9	2009-06-29 00:42:00 UTC	-73.986017	
199998	20259894	2015-05-20 14:56:25.0000004	14.5	2015-05-20 14:56:25 UTC	-73.997124	
199999	11951496	2010-05-15 04:08:00.00000076	14.1	2010-05-15 04:08:00 UTC	-73.984395	

200000 rows × 9 columns

```
In [32]: df.head()
```

Out[32]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	4
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	4
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	4
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	4
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	4

In [33]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #  Column            Non-Null Count  Dtype  
--- 
 0  Unnamed: 0        200000 non-null   int64  
 1  key               200000 non-null   object  
 2  fare_amount       200000 non-null   float64 
 3  pickup_datetime   200000 non-null   object  
 4  pickup_longitude  200000 non-null   float64 
 5  pickup_latitude   200000 non-null   float64 
 6  dropoff_longitude 199999 non-null   float64 
 7  dropoff_latitude  199999 non-null   float64 
 8  passenger_count   200000 non-null   int64  
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

In [34]: df.columns

```
Out[34]: Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
                 'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
                 'dropoff_latitude', 'passenger_count'],
                dtype='object')
```

In [35]: df = df.drop(['Unnamed: 0', 'key'], axis=1)

In [36]: df.head()

Out[36]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dr
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	
2	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	
3	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	
4	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	

C C

In [37]: df.shape

Out[37]: (200000, 7)

In [38]: df.dtypes

Out[38]:

fare_amount	float64
pickup_datetime	object
pickup_longitude	float64
pickup_latitude	float64
dropoff_longitude	float64
dropoff_latitude	float64
passenger_count	int64
dtype:	object

In [39]: df.describe()

Out[39]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000
mean	11.359955	-72.527638	39.935885	-72.525292	39.92389
std	9.901776	11.437787	7.720539	13.117408	6.79482
min	-52.000000	-1340.648410	-74.015515	-3356.666300	-881.98551
25%	6.000000	-73.992065	40.734796	-73.991407	40.73382
50%	8.500000	-73.981823	40.752592	-73.980093	40.75304
75%	12.500000	-73.967154	40.767158	-73.963658	40.76800
max	499.000000	57.418457	1644.421482	1153.572603	872.69762

C C

In [40]: df.isnull().sum()

```
Out[40]: fare_amount      0
          pickup_datetime    0
          pickup_longitude    0
          pickup_latitude     0
          dropoff_longitude   1
          dropoff_latitude    1
          passenger_count     0
          dtype: int64
```

```
In [41]: df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(), inplace=True)
```

```
In [42]: df.isnull().sum()
```

```
Out[42]: fare_amount      0
          pickup_datetime    0
          pickup_longitude    0
          pickup_latitude     0
          dropoff_longitude   1
          dropoff_latitude    0
          passenger_count     0
          dtype: int64
```

```
In [43]: df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(), inplace=True)
```

```
In [44]: df.isnull().sum()
```

```
Out[44]: fare_amount      0
          pickup_datetime    0
          pickup_longitude    0
          pickup_latitude     0
          dropoff_longitude   0
          dropoff_latitude    0
          passenger_count     0
          dtype: int64
```

```
In [45]: df.dtypes
```

```
Out[45]: fare_amount      float64
          pickup_datetime    object
          pickup_longitude    float64
          pickup_latitude     float64
          dropoff_longitude   float64
          dropoff_latitude    float64
          passenger_count     int64
          dtype: object
```

```
In [46]: df.pickup_datetime = pd.to_datetime(df.pickup_datetime, errors='coerce')
```

```
In [47]: df.dtypes
```

```
Out[47]: fare_amount           float64
pickup_datetime      datetime64[ns, UTC]
pickup_longitude        float64
pickup_latitude         float64
dropoff_longitude        float64
dropoff_latitude         float64
passenger_count          int64
dtype: object
```

```
In [48]: df= df.assign(hour = df.pickup_datetime.dt.hour,
                     day = df.pickup_datetime.dt.day,
                     month = df.pickup_datetime.dt.month,
                     year = df.pickup_datetime.dt.year,
                     dayofweek = df.pickup_datetime.dt.dayofweek)
```

```
In [49]: df.head
```

```
Out[49]: <bound method NDFrame.head of
p_longitude \>
0 7.5 2015-05-07 19:52:06+00:00 -73.999817
1 7.7 2009-07-17 20:04:56+00:00 -73.994355
2 12.9 2009-08-24 21:45:00+00:00 -74.005043
3 5.3 2009-06-26 08:22:21+00:00 -73.976124
4 16.0 2014-08-28 17:47:00+00:00 -73.925023
...
199995 ... ...
199996 3.0 2012-10-28 10:49:00+00:00 -73.987042
199997 7.5 2014-03-14 01:09:00+00:00 -73.984722
199998 30.9 2009-06-29 00:42:00+00:00 -73.986017
199999 14.5 2015-05-20 14:56:25+00:00 -73.997124
199999 14.1 2010-05-15 04:08:00+00:00 -73.984395

pickup_latitude dropoff_longitude dropoff_latitude passenger_count \
0 40.738354 -73.999512 40.723217 1
1 40.728225 -73.994710 40.750325 1
2 40.740770 -73.962565 40.772647 1
3 40.790844 -73.965316 40.803349 3
4 40.744085 -73.973082 40.761247 5
...
199995 ... ...
199996 40.736837 -74.006672 40.739620 1
199997 40.756487 -73.858957 40.692588 2
199998 40.725452 -73.983215 40.695415 1
199999 40.720077 -73.985508 40.768793 1

hour day month year dayofweek
0 19 7 5 2015 3
1 20 17 7 2009 4
2 21 24 8 2009 0
3 8 26 6 2009 4
4 17 28 8 2014 3
...
199995 ... ...
199996 10 28 10 2012 6
199997 1 14 3 2014 4
199998 0 29 6 2009 0
199999 14 20 5 2015 2
199999 4 15 5 2010 5

[200000 rows x 12 columns]>
```

```
In [50]: df = df.drop('pickup_datetime', axis=1)
```

```
In [51]: df.head
```

```
Out[51]: <bound method NDFrame.head of
          de dropoff_longitude  \
          0           7.5    -73.999817     40.738354      -73.999512
          1           7.7    -73.994355     40.728225      -73.994710
          2          12.9    -74.005043     40.740770      -73.962565
          3           5.3    -73.976124     40.790844      -73.965316
          4          16.0    -73.925023     40.744085      -73.973082
          ...
          ...       ...
          199995      3.0    -73.987042     40.739367      -73.986525
          199996      7.5    -73.984722     40.736837      -74.006672
          199997     30.9    -73.986017     40.756487      -73.858957
          199998     14.5    -73.997124     40.725452      -73.983215
          199999     14.1    -73.984395     40.720077      -73.985508

          dropoff_latitude  passenger_count  hour  day  month  year  dayofweek
          0           40.723217            1   19   7    5  2015      3
          1           40.750325            1   20  17    7  2009      4
          2           40.772647            1   21  24    8  2009      0
          3           40.803349            3   8   26    6  2009      4
          4           40.761247            5   17  28    8  2014      3
          ...
          ...       ...
          199995      40.740297            1   10  28   10  2012      6
          199996      40.739620            1    1  14    3  2014      4
          199997      40.692588            2    0  29    6  2009      0
          199998      40.695415            1   14  20    5  2015      2
          199999      40.768793            1    4  15    5  2010      5

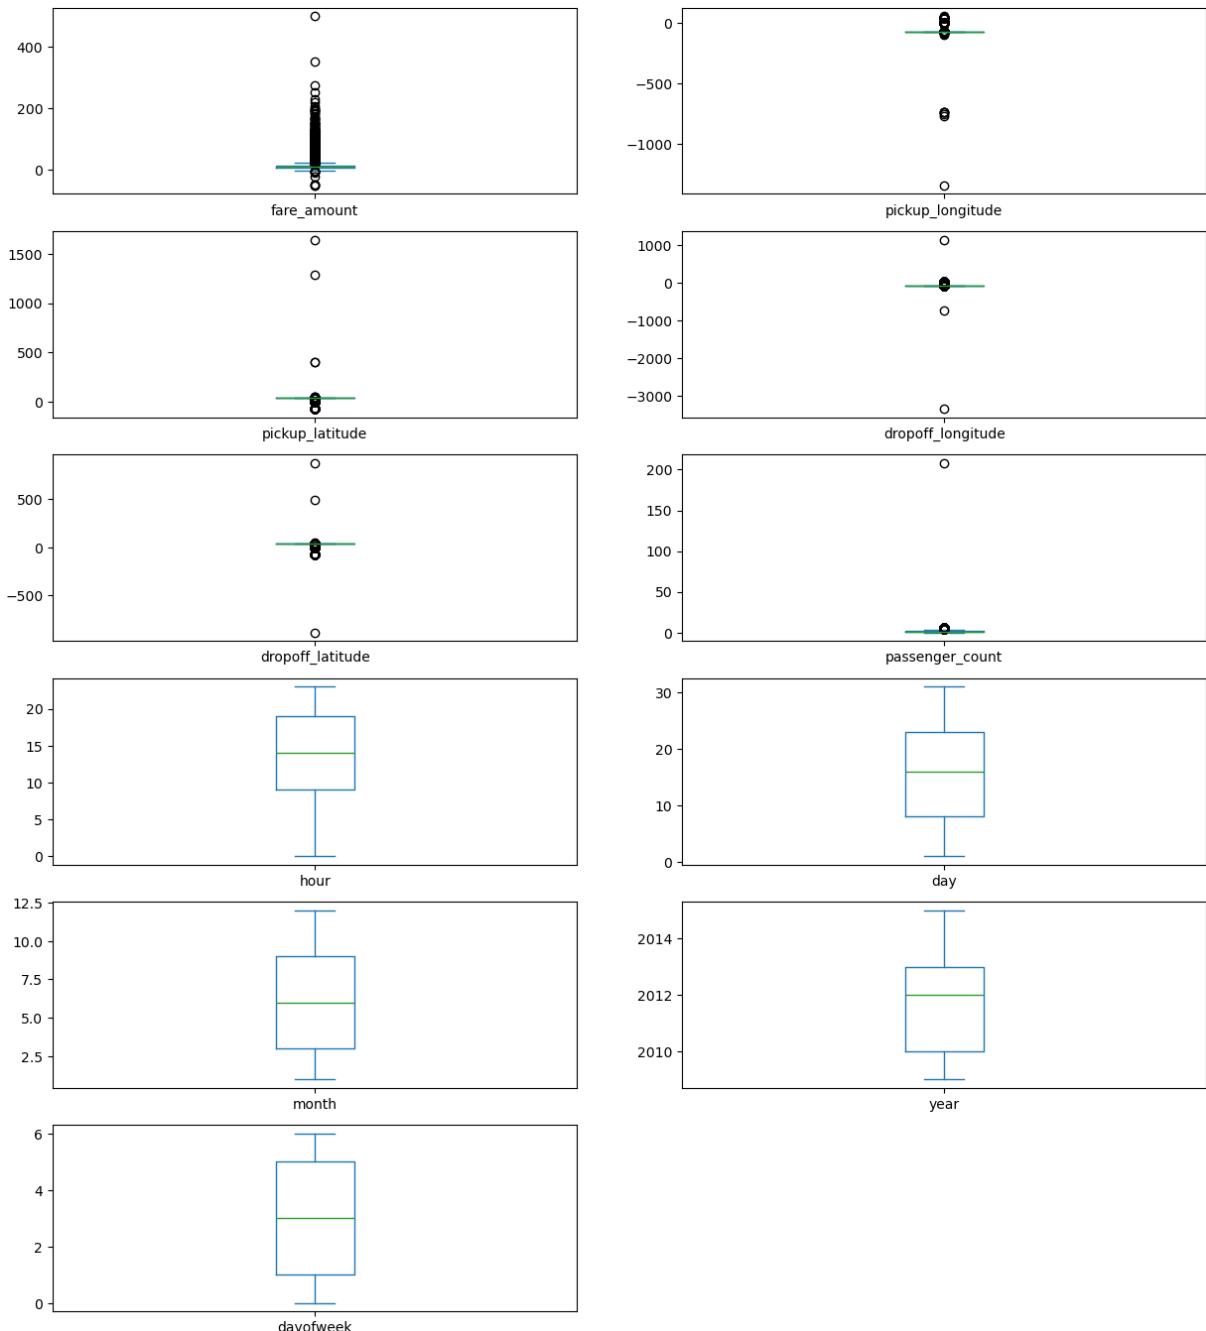
[ 200000 rows x 11 columns]>
```

```
In [52]: df.dtypes
```

```
Out[52]: fare_amount      float64
          pickup_longitude  float64
          pickup_latitude   float64
          dropoff_longitude float64
          dropoff_latitude  float64
          passenger_count   int64
          hour              int32
          day               int32
          month             int32
          year              int32
          dayofweek         int32
          dtype: object
```

```
In [53]: df.plot(kind="box", subplots = True, layout =(7,2), figsize=(15,20))
```

```
Out[53]: fare_amount      Axes(0.125,0.786098;0.352273x0.0939024)
pickup_longitude    Axes(0.547727,0.786098;0.352273x0.0939024)
pickup_latitude      Axes(0.125,0.673415;0.352273x0.0939024)
dropoff_longitude    Axes(0.547727,0.673415;0.352273x0.0939024)
dropoff_latitude      Axes(0.125,0.560732;0.352273x0.0939024)
passenger_count      Axes(0.547727,0.560732;0.352273x0.0939024)
hour                  Axes(0.125,0.448049;0.352273x0.0939024)
day                   Axes(0.547727,0.448049;0.352273x0.0939024)
month                 Axes(0.125,0.335366;0.352273x0.0939024)
year                  Axes(0.547727,0.335366;0.352273x0.0939024)
dayofweek             Axes(0.125,0.222683;0.352273x0.0939024)
dtype: object
```



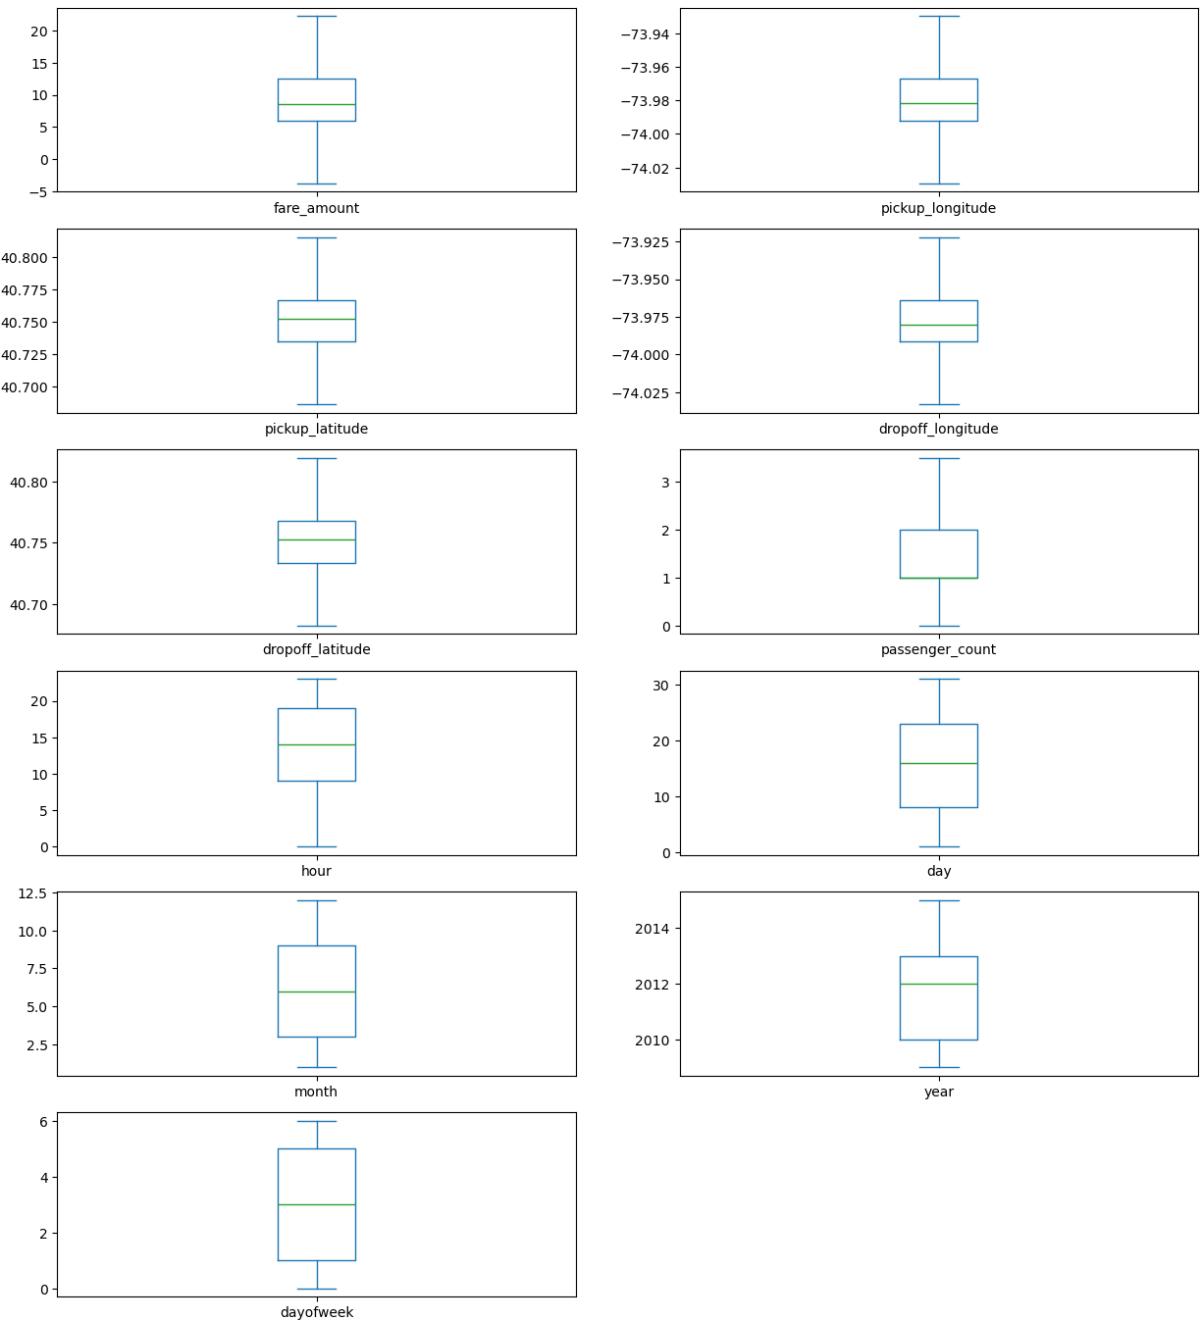
```
In [62]: def remove_outlier(df1, col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
```

```
IQR= Q3 - Q1
lower_whisker= Q1-1.5*IQR
upper_whisker= Q3+1.5*IQR
df[col] = np.clip(df1[col], lower_whisker, upper_whisker)
return df1
def treat_outliers_all(df1, col_list):
    for c in col_list:
        df1 = remove_outlier(df, c)
    return df1
```

```
In [63]: df = treat_outliers_all(df,df.iloc[:, 0::])
```

```
In [64]: df.plot(kind = "box", subplots = True, layout=(7,2), figsize=(15,20))
```

```
Out[64]: fare_amount      Axes(0.125,0.786098;0.352273x0.0939024)
pickup_longitude     Axes(0.547727,0.786098;0.352273x0.0939024)
pickup_latitude       Axes(0.125,0.673415;0.352273x0.0939024)
dropoff_longitude     Axes(0.547727,0.673415;0.352273x0.0939024)
dropoff_latitude       Axes(0.125,0.560732;0.352273x0.0939024)
passenger_count       Axes(0.547727,0.560732;0.352273x0.0939024)
hour                  Axes(0.125,0.448049;0.352273x0.0939024)
day                   Axes(0.547727,0.448049;0.352273x0.0939024)
month                 Axes(0.125,0.335366;0.352273x0.0939024)
year                  Axes(0.547727,0.335366;0.352273x0.0939024)
dayofweek              Axes(0.125,0.222683;0.352273x0.0939024)
dtype: object
```



```
In [65]: pip install haversine
```

```
Collecting haversine
  Downloading haversine-2.9.0-py2.py3-none-any.whl.metadata (5.8 kB)
  Downloading haversine-2.9.0-py2.py3-none-any.whl (7.7 kB)
Installing collected packages: haversine
Successfully installed haversine-2.9.0
Note: you may need to restart the kernel to use updated packages.
```

```
In [72]: import haversine as hs
travel_dist = []

for pos in range(len(df['pickup_longitude'])):
    long1 = df['pickup_longitude'][pos]
    lat1 = df['pickup_latitude'][pos]
    long2 = df['dropoff_longitude'][pos]
    lat2 = df['dropoff_latitude'][pos]
    travel_dist.append(hs.haversine((long1, lat1), (long2, lat2)))
```

```

lati2 = df['dropoff_latitude'][pos] # <-- corrected from 'dropoff_lloc1'

loc1 = (lati1, long1)
loc2 = (lati2, long2)

c = hs.haversine(loc1, loc2)
travel_dist.append(c)

df['dist_travel_km'] = travel_dist
df.head()

```

Out[72]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	pa
0	7.5	-73.999817	40.738354	-73.999512	40.723217	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	
4	16.0	-73.929786	40.744085	-73.973082	40.761247	

C C

In [73]:

```
df = df.loc[(df.dist_travel_km>=1) | (df.dist_travel_km<=130) ]
print("Remaining obervation: ", df.shape)
```

Remaining obervation: (200000, 12)

In [74]:

```
incorrect_coordinates = df.loc[(df.pickup_latitude>90) | (df.pickup_latitude<-90) |
                               (df.dropoff_latitude>90) | (df.dropoff_latitude<-90) |
                               (df.pickup_longitude>180) | (df.pickup_longitude<-180) |
                               (df.dropoff_longitude>90) | (df.dropoff_longitude<-90) ]
```

In [75]:

```
df.drop(incorrect_coordinates, inplace = True, errors ='ignore')
```

In [76]:

```
df.head()
```

Out[76]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	pa
0	7.5	-73.999817	40.738354	-73.999512	40.723217	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	
4	16.0	-73.929786	40.744085	-73.973082	40.761247	

C C

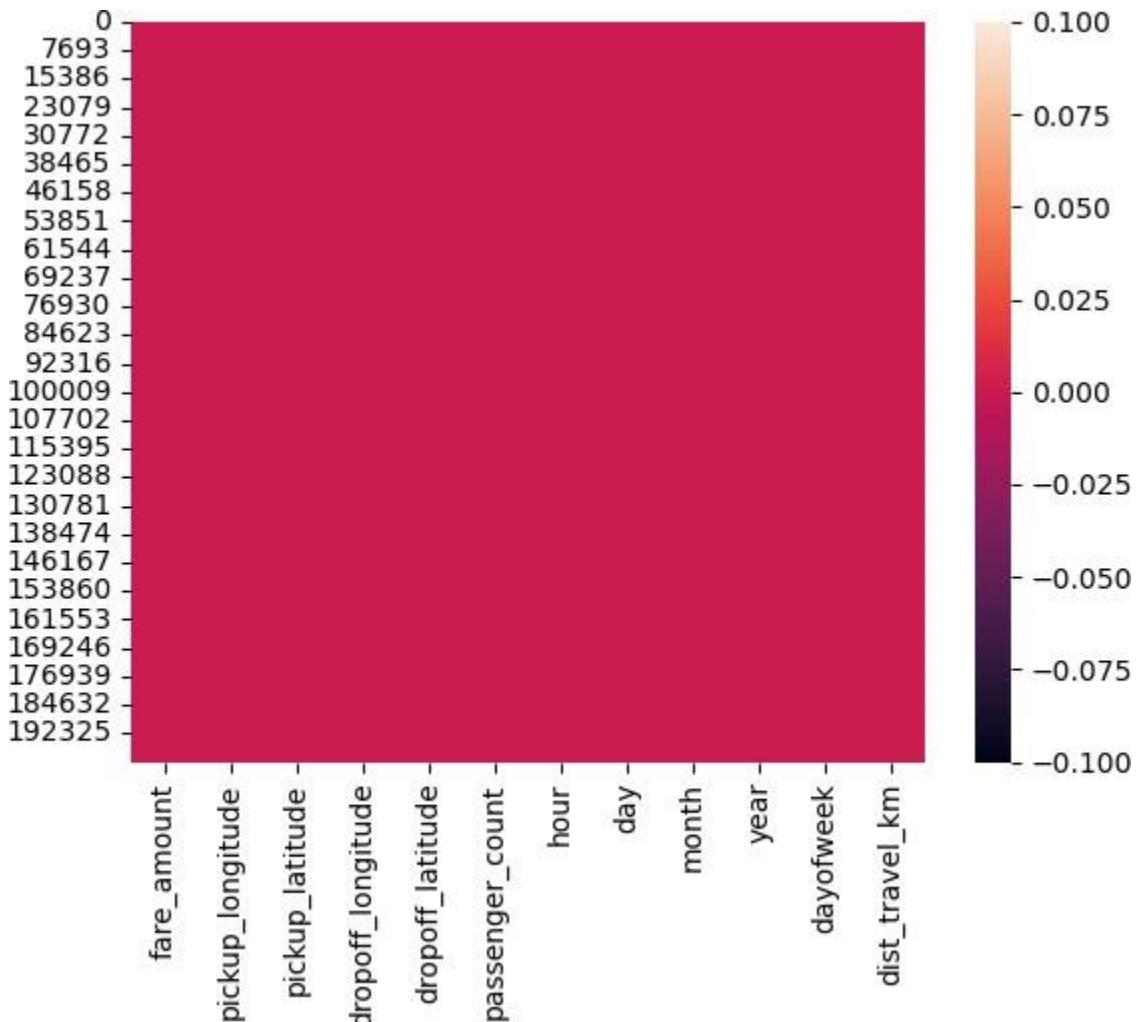
In [77]:

```
df.isnull().sum()
```

```
Out[77]: fare_amount      0  
pickup_longitude     0  
pickup_latitude       0  
dropoff_longitude     0  
dropoff_latitude       0  
passenger_count      0  
hour                  0  
day                   0  
month                 0  
year                  0  
dayofweek              0  
dist_travel_km        0  
dtype: int64
```

```
In [78]: sns.heatmap(df.isnull())
```

```
Out[78]: <Axes: >
```



```
In [79]: corr = df.corr()  
corr
```

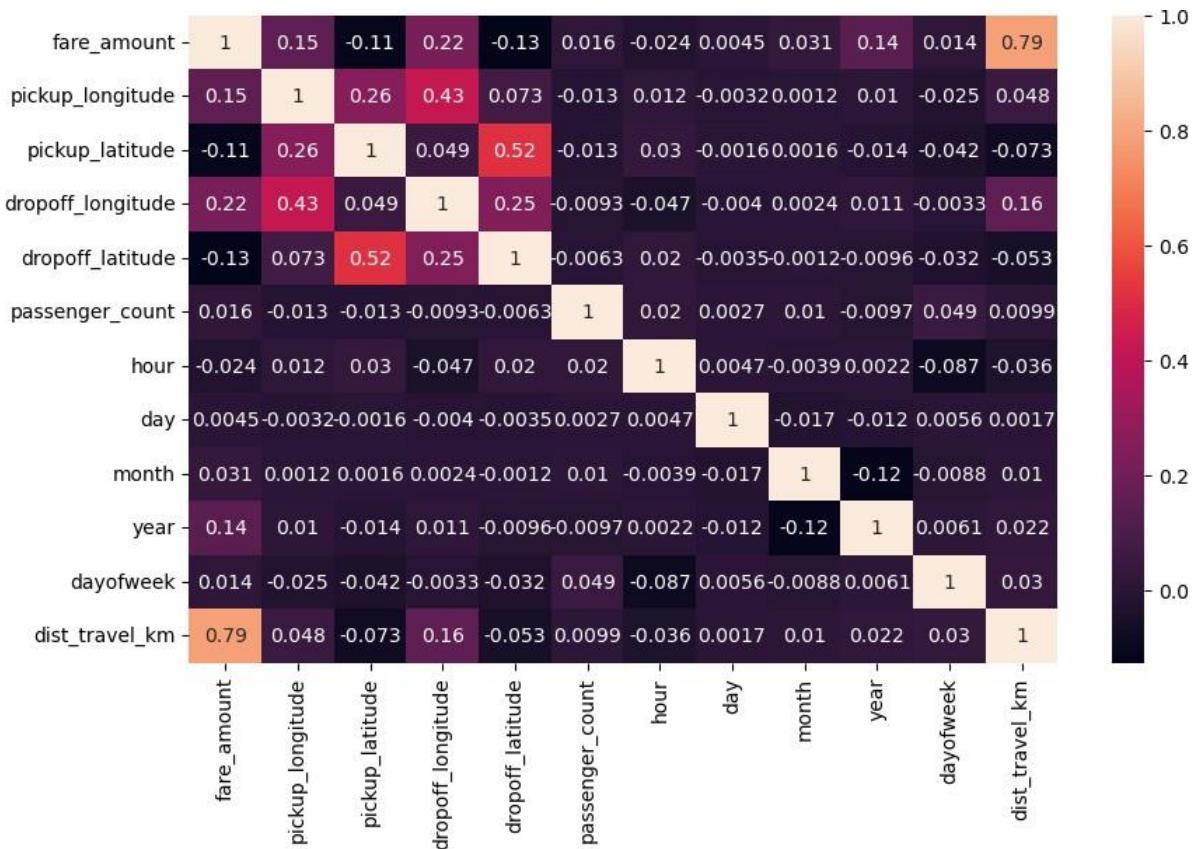
Out[79]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	drop
fare_amount	1.000000	0.154069	-0.110842	0.218675	
pickup_longitude	0.154069	1.000000	0.259497	0.425619	
pickup_latitude	-0.110842	0.259497	1.000000	0.048889	
dropoff_longitude	0.218675	0.425619	0.048889	1.000000	
dropoff_latitude	-0.125898	0.073290	0.515714	0.245667	
passenger_count	0.015778	-0.013213	-0.012889	-0.009303	
hour	-0.023623	0.011579	0.029681	-0.046558	
day	0.004534	-0.003204	-0.001553	-0.004007	
month	0.030817	0.001169	0.001562	0.002391	
year	0.141277	0.010198	-0.014243	0.011346	
dayofweek	0.013652	-0.024652	-0.042310	-0.003336	
dist_travel_km	0.786385	0.048446	-0.073362	0.155191	

C C

In [80]: `fig, axis = plt.subplots(figsize= (10,6))
sns.heatmap(df.corr(), annot = True)`

Out[80]: <Axes: >



```
In [81]: df.dtypes
```

```
Out[81]: fare_amount           float64
pickup_longitude            float64
pickup_latitude             float64
dropoff_longitude            float64
dropoff_latitude             float64
passenger_count              float64
hour                         int32
day                          int32
month                        int32
year                         int32
dayofweek                     int32
dist_travel_km                float64
dtype: object
```

```
In [84]: x = df[['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'passenger_count']]
```

```
In [85]: y = df['fare_amount']
```

```
In [87]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.33)
```

```
In [89]: from sklearn.linear_model import LinearRegression
regression = LinearRegression()
```

```
In [90]: regression.fit(x_train, y_train)
```

```
Out[90]: ▾ LinearReg ession ⓘ ?  
LinearRegression()
```

```
In [91]: regression.intercept_
```

```
Out[91]: 3256.513201301484
```

```
In [92]: regression.coef_
```

```
Out[92]: array([ 3.10758620e+01, -1.85910612e+01,   1.27429851e+01,   5.41800674e-02,
      5.77478080e-03,   3.38249411e-03,   5.89651957e-02,   3.71366494e-01,
     -3.27278737e-02,   1.86084919e+00])
```

```
In [93]: prediction = regression.predict(x_test)
```

```
In [94]: print(prediction)
```

```
[13.00853888  7.57316671  13.25903141 ... 13.71316448 13.18778424
 11.11547073]
```

```
In [95]: y_test
```

```
Out[95]:    1086      22.25
             96074     11.30
            104027     12.50
            133659      7.50
            58255      7.70
            ...
           27046      10.50
          178816     14.00
          42112      12.50
          53060      14.50
          83308      8.90
Name: fare_amount, Length: 66000, dtype: float64
```

```
In [97]: from sklearn.metrics import r2_score
r2_score(y_test,prediction)
```

```
Out[97]: 0.6584561379465177
```

```
In [100...]: from sklearn.metrics import mean_squared_error
MSE = mean_squared_error(y_test,prediction)
MSE
```

```
Out[100...]: 10.140972773285915
```

```
In [ ]:
```